

---

## 1. The problem

We have a dataset with **two features** — let's call them ( $x_1$ ) and ( $x_2$ ).

For example, these could represent:

- ( $x_1$ ): score on test 1
- ( $x_2$ ): score on test 2

If we train **logistic regression** directly on these two features, the model will create a **linear decision boundary** (a straight line).

But sometimes, the data is **not linearly separable** — meaning a straight line cannot divide the two classes correctly.

---

## 2. Feature mapping (creating more features)

To make the model more flexible, we **create new features** from the existing ones — this is called **feature mapping** or **polynomial feature expansion**.

Example:

Instead of only using  $x_1$  and  $x_2$ , we create terms like:

$$x_1^2, x_2^2, x_1x_2, x_1^3, x_1^2x_2, \dots$$

Up to the **6th power**, this creates **28 total features** (including all combinations of ( $x_1$ ) and ( $x_2$ )).

So our original 2D input (( $x_1$ ,  $x_2$ )) becomes a **28-dimensional vector** after mapping.

---

## 3. Why do this?

By training logistic regression on this **higher-dimensional data**, the model can draw **nonlinear decision boundaries** in the original 2D space.

So instead of a straight line, it might form a curved boundary that fits the data better.

---

## 4. The risk: Overfitting

While this gives the model **more power** (it can fit complex shapes), it also means the model can **overfit** — that is, memorize the training data instead of learning general patterns.

Overfitting happens when:

- The model performs very well on training data,

- But poorly on unseen (test) data.
- 

## 5. The solution: Regularization

To prevent overfitting, we use **regularization** — a technique that **penalizes large parameter values** (large coefficients).

This keeps the model simpler and helps it generalize better.

---

### Summary

Concept	Explanation
<b>Feature Mapping</b>	Expands original features into polynomial combinations (up to degree 6 here).
<b>Effect</b>	Makes logistic regression capable of fitting <b>nonlinear</b> decision boundaries.
<b>Resulting Dimension</b>	2 input features → 28 mapped features.
<b>Risk</b>	Model becomes complex and may <b>overfit</b> .
<b>Fix</b>	Use <b>regularization</b> to control complexity.

---

### Intuition — how we *penalize* large parameters

Logistic regression learns parameters (weights)  $\theta = [\theta_0, \theta_1, \dots, \theta_n]$ . Without regularization the cost (loss) is:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[ y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)})) \right]$$

where  $h_\theta(x) = \sigma(\theta^\top x)$  is the sigmoid prediction.

**Regularization** adds a penalty term to the cost that grows with the square of the parameters:

$$J_{\text{reg}}(\theta) = J(\theta) + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Key points:

- $\lambda$  (lambda) is the regularization parameter. Bigger  $\lambda \rightarrow$  heavier penalty.
- We usually **do not** penalize  $\theta_0$  (the bias/intercept); only the other parameters  $\theta_1..,\theta_n$ .
- The penalty forces the learning algorithm to prefer solutions with **smaller weights** (smaller magnitude of  $\theta_j$ ). Small weights correspond to simpler models (less wiggly / less likely to overfit).

## How it changes training (gradient)

The gradient used to update parameters also changes. For  $j = 0$  (bias):

$$\frac{\partial J_{\text{reg}}}{\partial \theta_0} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

For  $j \geq 1$ :

$$\frac{\partial J_{\text{reg}}}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j$$

So during gradient descent you update  $\theta_j$  using the extra  $\frac{\lambda}{m} \theta_j$  term — that term *pushes*  $\theta_j$  toward zero each step (shrinking the weights).

---

## What regularization does in practice

- **No regularization ( $\lambda = 0$ )**: the model can make very large weights to fit the training data exactly — often leads to overfitting and a complex decision boundary.
- **Moderate regularization (small  $\lambda$ )**: reduces overfitting while keeping enough flexibility to fit true patterns.
- **Large regularization (big  $\lambda$ )**: forces weights to be very small; the model becomes simple and may underfit (too rigid).

So the aim is to choose ( $\lambda$ ) so the model generalizes well (often via cross-validation).

---



### 1. What are the “weights” ( $\theta$ 's)?

In logistic regression, we predict the probability of belonging to class 1 as

$$h_{\theta}(x) = \sigma(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n)$$

where each  $\theta_i$  tells us *how strongly* its feature ( $x_i$ ) affects the output.

- If  $\theta_i$  is large (positive) → that feature pushes the prediction strongly toward class 1.
- If  $\theta_i$  is large (negative) → that feature pushes the prediction strongly toward class 0.
- If  $\theta_i \approx 0$  → that feature has little effect on the prediction.

So the **magnitude**  $|\theta_i|$  measures how “powerful” that feature’s influence is.

---



## 2. What “penalizing large weights” means

Regularization adds a term to the cost function:

$$\frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

This term **grows quickly when any  $\theta_i$  becomes large**, because we square it.

During training, the optimizer tries to **minimize the total cost** — and the only way to keep this penalty small is to **keep  $\theta_i$  small**.

👉 That's why we say regularization “**penalizes large weights**.”

---



## 3. Why we want small weights

When the weights are **very large**, the model’s output (after the sigmoid) changes **very abruptly** with small variations in the inputs — making the decision boundary extremely wiggly and sensitive to noise.

When the weights are **smaller**, the model’s output changes more smoothly — it captures only the broad, important trend instead of memorizing every fluctuation.

Effect	Large weights	Small weights
Sensitivity to input	Very high (unstable)	Moderate (stable)
Decision boundary	Complex, wiggly	Smooth, general
Risk of overfitting	High	Low
Model interpretability	Hard	Easier

---



## 4. Visual intuition

Imagine fitting a curve to noisy data points.

- **Without regularization (large  $\theta$ 's allowed):**  
The curve bends a lot to go through every point — fits noise → overfitting.
  - **With regularization (small  $\theta$ 's forced):**  
The curve stays smoother — ignores small fluctuations → better generalization.
-

## 5. Balance: not too large, not too small

- If  $\lambda = 0$ , no penalty  $\rightarrow \theta$ 's can explode  $\rightarrow$  overfitting.
  - If  $\lambda$  is too big, all  $\theta$ 's shrink nearly to 0  $\rightarrow$  model becomes too simple (underfitting).
  - The goal is to find a **sweet spot** where  $\theta$ 's are small enough to be stable, but not so small that the model loses its expressive power.
- 

### In short:

**Penalizing large weights** means forcing the model's parameters ( $\theta$ 's) to stay small so the model makes smoother, more general predictions instead of memorizing the training data.

---